

BECOME AN AI ENGINEER — NOT JUST AN AI USER

Agentic AI Engineering Cohort

Stop watching AI tutorials. Start building AI products.

This premium 10-week Agentic AI Engineering Bootcamp is built for working software engineers & students who want practical, production-ready skills. You'll design and build RAG systems, tool-calling workflows, MCP servers, and autonomous AI agents while developing a single real-world capstone project that grows every week—from the first coding session to production deployment.

10

weeks, hands-on

Sat or Sunweekends only, 4–5
hrs/day**Pune**

offline + online access

≤ 25

students per cohort

1deployed capstone
product

WHAT YOU'LL BUILD

- LLM apps with streaming responses

- RAG systems on vector databases

- Tool-using autonomous systems

- AI agents & multi-agent workflows

- Context-aware AI assistants with memory & actions

- Production AI apps, deployed live

Every week ships something real.



Saturday/Sunday · 4-5 hrs/day · one evolving capstone from week 1 → week 10

WEEK

01.02

SHIPS ▶

Terminal chatbot
with structured
outputs

AI Foundations & Prompt Engineering

Understand how LLMs work under the hood — and drive them with precision.

- LLM internals & how GPT works
- Tokens, context windows & cost
- Structured outputs & Pydantic
- First API call — terminal chatbot
- Prompt engineering fundamentals
- Temperature, top-p & sampling

Python 3.12

OpenAI SDK

Pydantic

WEEK

03

SHIPS ▶

“Ask my notes”
semantic search
app

Context Engineering & Embeddings

Manage memory and make machines understand meaning, not keywords.

- Context windows & prompt packing
- What embeddings really are
- Similarity scoring & top-k
- Sliding-window memory
- Semantic vs keyword search
- Build: “Ask my notes”

Embeddings

Semantic search

NumPy

WEEK

04

SHIPS ▶

PDF chatbot with
citations

Vector Databases & RAG

Build production-grade retrieval over your own documents.

- ChromaDB setup & queries
- Dense retrieval & metadata filters
- PDF chatbot with citations
- Chunking strategies
- RAG architecture, end-to-end
- Multi-document querying

ChromaDB

PyPDF

RAG

WEEK

05

SHIPS ▶

Tool-calling
orchestrator

Advanced RAG & Tool Calling

Sharpen retrieval and let the model take real actions in the world.

- Query rewriting & hybrid search
- Build tools: calculator, file reader
- AI as an orchestrator
- Re-ranking & context compression
- Function execution loop
- Error & retry handling

Function calling

Hybrid search

Re-ranking

From agents to production deploy.

WEEKS 06-10 • CONTINUED

WEEK 06

SHIPS ▶
Autonomous
research agent

AI Agents

Give your AI autonomy — planning, memory and self-correction.

- Agent lifecycle & reasoning loops
- Short & long-term agent memory
- Tool selection & routing
- Planning, reflection & retry
- Task decomposition
- Build: research agent

Agent loops

Memory stores

ReAct

WEEK 07

SHIPS ▶
MCP-powered
filesystem
assistant

Multi-Agent Systems & MCP

Coordinate teams of agents and connect them to real systems.

- Supervisor-worker pattern
- What MCP is & why it matters
- Agent-to-agent messaging
- Researcher + summarizer agents
- Set up an MCP server
- AI filesystem assistant

LangGraph

MCP

Multi-agent

WEEK 08

SHIPS ▶
Real-time voice
agent +
hardened API

Voice Agents & Production Engineering

Give your AI a real-time voice — then make it fast, cheap and safe.

- Voice agent architecture (STT → LLM → TTS)
- Interruptions & low-latency responses
- Latency & cost optimization
- Prompt injection & guardrails
- Real-time streaming audio & turn-taking
- Build: voice assistant with tool calls
- Streaming APIs & rate limiting
- Caching & token budgeting

Speech-to-Text

Text-to-Speech

Realtime voice

Streaming APIs

Guardrails

WEEK 09

SHIPS ▶
Live deploy +
eval suite

Evals & Deployment

Measure quality objectively and ship your system to a public URL.

- Lightweight eval suite
- FastAPI deployment
- Secrets & env management
- Prompt regression testing
- Deploy to Railway / Render
- CI/CD basics

FastAPI

Railway / Render

CI/CD

WEEK 10

SHIPS ▶
Deployed AI
Engineering
Assistant

Final Capstone & Demo Day

Integrate everything into one deployed product and present it live.

- End-to-end system integration
- Final live demos
- Portfolio & code polish
- Architecture review
- Peer reviews
- Career roadmap discussion

End-to-end

Demo Day

“AI Engineering Assistant” — your deployed product.



You don't graduate with a certificate — you graduate with a live, working AI product on a public URL, built progressively across all ten weeks.

Core AI Knowledge

- Chat interface
- Streaming AI responses
- System prompts & context handling
- Conversation memory
- Structured JSON outputs

Agentic Engineering

- PDF ingestion & semantic search
- Vector database + RAG pipeline
- Multi-document, citation-aware answers
- Task planning & tool-using agents
- Multi-agent + MCP integrations

Production

- Logging & token tracking
- Lightweight evaluations
- Prompt versioning
- Error handling
- Deployed publicly · portfolio-ready

WHO IT'S FOR

Software Engineers

Developers with 2–5 years' experience moving into AI engineering and modern AI application development.

Backend & Full-Stack

Engineers building intelligent systems, APIs and scalable AI products on top of existing skills.

Professionals Exploring AI

People already using AI tools who now want to build and deploy AI systems professionally.

THE STACK YOU'LL USE

Python 3.12

FastAPI

OpenAI SDK

Claude

Gemini

LangChain

LangGraph

ChromaDB

FAISS

pgvector

MCP

Streamlit

React

Docker

Railway

Render

WHY REBINDRISE

- Industry-aligned, code-first curriculum
- 100% in-person, hands-on training
- Real projects you can show off
- Expert mentors — 10+ years in tech
- Small batches, personal attention

Designed for serious software engineers.

No hand-holding — real engineering depth, from week one.

Office No. 604, 6th Floor, Suratwala Mark Plazzo, Hinjawadi, Pune, Maharashtra — 411057

rebindrise.com